

```
In [1]: # William Barker
# DSC680
# Project 3

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import IsolationForest
from sklearn.metrics import confusion_matrix, classification_report, roc_auc_score, roc_curve

# Set random seed for reproducibility
np.random.seed(42)
```

```
In [3]: # Load the dataset
df = pd.read_csv('creditcard_2023.csv')

# Check the shape and first few rows of the dataset
print(df.shape)
print(df.head())
```

```
(568630, 31)
   id  V1      V2      V3      V4      V5      V6      V7  \
0    0 -0.260648 -0.469648  2.496266 -0.083724  0.129681  0.732898  0.519014
1    1  0.985100 -0.356045  0.558056 -0.429654  0.277140  0.428605  0.406466
2    2 -0.260272 -0.949385  1.728538 -0.457986  0.074062  1.419481  0.743511
3    3 -0.152152 -0.508959  1.746840 -1.090178  0.249486  1.143312  0.518269
4    4 -0.206820 -0.165280  1.527053 -0.448293  0.106125  0.530549  0.658849

      V8      V9  ...      V21      V22      V23      V24      V25  \
0 -0.130006  0.727159  ... -0.110552  0.217606 -0.134794  0.165959  0.126280
1 -0.133118  0.347452  ... -0.194936 -0.605761  0.079469 -0.577395  0.190090
2 -0.095576 -0.261297  ... -0.005020  0.702906  0.945045 -1.154666 -0.605564
3 -0.065130 -0.205698  ... -0.146927 -0.038212 -0.214048 -1.893131  1.003963
4 -0.212660  1.049921  ... -0.106984  0.729727 -0.161666  0.312561 -0.414116

      V26      V27      V28  Amount  Class
0 -0.434824 -0.081230 -0.151045  17982.10    0
1  0.296503 -0.248052 -0.064512   6531.37    0
2 -0.312895 -0.300258 -0.244718   2513.54    0
3 -0.515950 -0.165316  0.048424   5384.44    0
4  1.071126  0.023712  0.419117  14278.97    0

[5 rows x 31 columns]
```

```
In [4]: # Summary statistics
print(df.describe())
```

```
count      id      V1      V2      V3      V4  \
count  568630.000000  5.686300e+05  5.686300e+05  5.686300e+05  5.686300e+05
mean    284314.500000 -1.109271e-14 -3.429498e-14 -1.209242e-14  3.825991e-15
std     164149.486121  1.000001e+00  1.000001e+00  1.000001e+00  1.000001e+00
min         0.000000 -3.495584e+00 -4.996657e+01 -3.183760e+00 -4.951222e+00
25%     142157.250000 -5.652859e-01 -4.866777e-01 -6.492987e-01 -6.560203e-01
50%     284314.500000 -9.363846e-02 -1.358939e-01  3.528579e-04 -7.376152e-02
75%     426471.750000  8.326582e-01  3.435552e-01  6.285380e-01  7.070047e-01
max     568629.000000  2.229046e+00  4.361865e+00  1.412583e+01  3.201536e+00
```

	V5	V6	V7	V8	V9 \
count	5.686300e+05	5.686300e+05	5.686300e+05	5.686300e+05	5.686300e+05
mean	6.288281e-15	-2.751174e-14	1.240002e-14	8.208047e-15	-1.002980e-14
std	1.000001e+00	1.000001e+00	1.000001e+00	1.000001e+00	1.000001e+00
min	-9.952786e+00	-2.111111e+01	-4.351839e+00	-1.075634e+01	-3.751919e+00
25%	-2.934955e-01	-4.458712e-01	-2.835329e-01	-1.922572e-01	-5.687446e-01
50%	8.108788e-02	7.871758e-02	2.333659e-01	-1.145242e-01	9.252647e-02
75%	4.397368e-01	4.977881e-01	5.259548e-01	4.729905e-02	5.592621e-01
max	4.271689e+01	2.616840e+01	2.178730e+02	5.958040e+00	2.027006e+01

	...	V21	V22	V23	V24 \
count	...	5.686300e+05	5.686300e+05	5.686300e+05	5.686300e+05
mean	...	2.210679e-15	-8.767441e-16	4.376179e-16	6.825608e-16
std	...	1.000001e+00	1.000001e+00	1.000001e+00	1.000001e+00
min	...	-1.938252e+01	-7.734798e+00	-3.029545e+01	-4.067968e+00
25%	...	-1.664408e-01	-4.904892e-01	-2.376289e-01	-6.515801e-01
50%	...	-3.743065e-02	-2.732881e-02	-5.968903e-02	1.590123e-02
75%	...	1.479787e-01	4.638817e-01	1.557153e-01	7.007374e-01
max	...	8.087080e+00	1.263251e+01	3.170763e+01	1.296564e+01

	V25	V26	V27	V28	Amount \
count	5.686300e+05	5.686300e+05	5.686300e+05	5.686300e+05	568630.000000
mean	2.545689e-15	1.781906e-15	2.817586e-15	2.891419e-15	12041.957635
std	1.000001e+00	1.000001e+00	1.000001e+00	1.000001e+00	6919.644449
min	-1.361263e+01	-8.226969e+00	-1.049863e+01	-3.903524e+01	50.010000
25%	-5.541485e-01	-6.318948e-01	-3.049607e-01	-2.318783e-01	6054.892500
50%	-8.193162e-03	-1.189208e-02	-1.729111e-01	-1.392973e-02	12030.150000
75%	5.500147e-01	6.728879e-01	3.340230e-01	4.095903e-01	18036.330000
max	1.462151e+01	5.623285e+00	1.132311e+02	7.725594e+01	24039.930000

	Class
count	568630.0
mean	0.5
std	0.5
min	0.0
25%	0.0
50%	0.5
75%	1.0
max	1.0

[8 rows x 31 columns]

In [5]:

```
# Check for missing values
print(df.isnull().sum())
```


```
id      0
V1      0
V2      0
V3      0
V4      0
V5      0
V6      0
V7      0
V8      0
V9      0
V10     0
V11     0
V12     0
V13     0
V14     0
V15     0
V16     0
V17     0
V18     0
V19     0
```

```
V20      0
V21      0
V22      0
V23      0
V24      0
V25      0
V26      0
V27      0
V28      0
Amount    0
Class     0
dtype: int64
```

```
In [6]: # Check the distribution of the 'Class' column
print(df['Class'].value_counts())
```


```
0      284315
1      284315
Name: Class, dtype: int64
```

```
In [7]: # Visualize class distribution
sns.countplot(x='Class', data=df)
plt.title('Class Distribution')
plt.show()
```



```
In [8]: # Correlation matrix
corr_matrix = df.corr()

# Visualize the correlation matrix
plt.figure(figsize=(12, 10))
sns.heatmap(corr_matrix, cmap='coolwarm', annot=False, fmt='.2f')
plt.title('Correlation Matrix')
plt.show()
```



```
In [9]: # Separate features and target variable
X = df.drop(columns=['id', 'Class'])
y = df['Class']

# Train-test split (80% training, 20% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42,

# Feature scaling (only for Logistic Regression)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [10]: # Initialize Logistic Regression model
log_reg = LogisticRegression(class_weight='balanced', random_state=42)

# Train the model
log_reg.fit(X_train_scaled, y_train)
```

```
Out[10]: LogisticRegression
LogisticRegression(class_weight='balanced', random_state=42)
```

```
In [11]: # Make predictions
y_pred_log_reg = log_reg.predict(X_test_scaled)
y_prob_log_reg = log_reg.predict_proba(X_test_scaled)[: , 1]


# Evaluation
print("Logistic Regression Classification Report:")
print(classification_report(y_test, y_pred_log_reg))
```

```
Logistic Regression Classification Report:
              precision    recall  f1-score   support

    0               0.95        0.98        0.97        56863
    1               0.98        0.95        0.96        56863

 accuracy               0.96        113726
 macro avg              0.97        0.96        0.96        113726
 weighted avg           0.97        0.96        0.96        113726
```


```
In [12]: # Confusion Matrix
cm_log_reg = confusion_matrix(y_test, y_pred_log_reg)
sns.heatmap(cm_log_reg, annot=True, fmt='d', cmap='Blues')
plt.title('Logistic Regression Confusion Matrix')
plt.show()
```



```
In [13]: # ROC-AUC Score
roc_auc_log_reg = roc_auc_score(y_test, y_prob_log_reg)
print(f'ROC-AUC Score for Logistic Regression: {roc_auc_log_reg:.4f}')
```

```
ROC-AUC Score for Logistic Regression: 0.9935
```

```
In [14]: # Plot ROC curve
fpr, tpr, thresholds = roc_curve(y_test, y_prob_log_reg)
plt.plot(fpr, tpr, label=f'Logistic Regression (AUC = {roc_auc_log_reg:.4f})')
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.show()
```



```
In [15]: # Initialize Isolation Forest model
iso_forest = IsolationForest(contamination=0.01, random_state=42)
```

```
In [16]: # Train the model (use training data only)
iso_forest.fit(X_train)
```

/Users/cameronbarker/opt/anaconda3/lib/python3.9/site-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but IsolationForest was fitted with feature names

```
warnings.warn(
```

```
Out[16]: ▼ IsolationForest
IsolationForest(contamination=0.01, random_state=42)
```

```
In [17]: # Predict anomalies (Isolation Forest predicts -1 for anomalies, 1 for normal)
y_pred_iso_forest = iso_forest.predict(X_test)
```

```
In [18]: # Convert predictions to match the binary classification format (0: normal, 1: fraud)
y_pred_iso_forest = np.where(y_pred_iso_forest == -1, 1, 0)
```

```
In [19]: # Evaluation
print("Isolation Forest Classification Report:")
print(classification_report(y_test, y_pred_iso_forest))
```

```
Isolation Forest Classification Report:
              precision    recall  f1-score   support


     0       0.50         1.00      0.67       56863
     1       0.97         0.02      0.04       56863

 accuracy          0.51       113726
 macro avg         0.74         0.51      0.35       113726
 weighted avg         0.74         0.51      0.35       113726
```

```
In [20]: # Confusion Matrix
cm_iso_forest = confusion_matrix(y_test, y_pred_iso_forest)
sns.heatmap(cm_iso_forest, annot=True, fmt='d', cmap='Blues')
plt.title('Isolation Forest Confusion Matrix')
```

```
plt.show()
```


```
# Since Isolation Forest doesn't give probability estimates, ROC-AUC isn't applicable
```



In [21]:

```
# Get the anomaly scores for the test set
anomaly_scores = iso_forest.decision_function(X_test)

# Higher scores correspond to more "normal" data points, while lower scores are more "anomalous"
plt.hist(anomaly_scores, bins=50)
plt.xlabel('Anomaly Score')
plt.ylabel('Frequency')
plt.title('Distribution of Anomaly Scores')
plt.show()
```



In []: